

HW8 Fully-Connected Neural Networks and "Dropout"

M113040105 劉東霖

壹. 所使用的 function:

一. rms_prop:

更新的公式如下圖所示:

$$\begin{aligned} &\textbf{RMSProp} \\ &v_t = \beta * v_{t-1} + (1 - \beta)(\nabla w_t)^2 \\ &w_{t+1} = w_t - \frac{\eta}{\sqrt{(v_t)} + \epsilon} \nabla w_t \end{aligned}$$

如下圖所示，我先把會用到的東西從 config 字典裡面拿出來:

```
grad_squared=config['cache']
learning_rate=config['learning_rate']
decay_rate=config['decay_rate']
epsilon=config['epsilon']
```

再來算出目前為止 gradient 的總和，這裡算總和的方式為 Exponentially weighted averages，因為 gradient 有正有負，所以一開始先平方相加，

要使用時再開根號即可。公式如下:

$$\begin{aligned} V_t &= \beta * V_{t-1} + (1 - \beta) * \theta_t \\ V_t &\text{ as approximately average} \\ &\text{over } \approx \frac{1}{(1-\beta)} \text{ days' temperature} \end{aligned}$$

程式碼如下:

```
grad_squared=decay_rate*grad_squared+(1-decay_rate)*dw*dw
```

再來就是更新梯度的部分。跟 sgd momentum 的差別在於多除了一項 $\sqrt{\text{grad_squared} + \text{epsilon}}$ 來調整學習率，epsilon 是為了防止 overflow。程式碼如下：

```
next_w=w-(learning_rate*dw)/(torch.sqrt(grad_squared)+epsilon)
```

最後再把所有東西存回字典裡面，程式碼如下：

```
config['cache']=grad_squared  
config['learning_rate']=learning_rate  
config['decay_rate']=decay_rate  
config['epsilon']=epsilon
```

二. adam:

更新的公式如下圖所示：

$$\begin{aligned} m_t &= \beta_1 * m_{t-1} + (1 - \beta_1) * \nabla w_t \\ v_t &= \beta_2 * v_{t-1} + (1 - \beta_2) * (\nabla w_t)^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t} & \hat{v}_t &= \frac{v_t}{1 - \beta_2^t} \\ w_{t+1} &= w_t - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} * \hat{m}_t \end{aligned}$$

如下圖所示，我先把會用到的東西從 config 字典裡面拿出來，然後 t 執行前先+1：

```
learning_rate=config['learning_rate']
epsilon=config['epsilon']
b1=config['beta1']
b2=config['beta2']
m=config['m']
v=config['v']
t=config['t']
t+=1
```

再來算出目前為止 Mean Squared Gradient 和 Root Mean Squared Gradient 的總和。程式碼如下：

```
m=b1*m+(1-b1)*dw
v=b2*v+(1-b2)*dw**2
```

再來修正 Mean Squared Gradient 和 Root Mean Squared Gradient 的偏差，因 gradient 的總和可能會受到過大或過小的梯度估計值的影響，進而導致收斂速度變慢或不穩定。因此，修正偏差估計是 Adam 算法的一個重要步驟，有助於改進其性能。程式碼如下：

```
m_unbias=m/(1-b1**t)
v_unbias=v/(1-b2**t)
```

再來就是更新梯度的部分，程式碼如下：

```
next_w=w-learning_rate*m_unbias/(torch.sqrt(v_unbias)+epsilon)
```

最後再把所有東西存回字典裡面，程式碼如下：

```
config['m']=m
config['v']=v
config['t']=t
config['learning_rate']=learning_rate
config['epsilon']=epsilon
config['beta1']=b1
config['beta2']=b2
```

三.dropout forward and backward:

Dropout forward 在訓練的時候，會隨機生成 0 到 1 的數，假設 $p=0.4$ ，代表 40% 會被 dropout，60% 會保持原本的輸出。除以 p 的部分是讓 train 和 test 的值一樣。在測試的時候不需要 dropout，值直接 pass 過去。程式碼如下：

```
if mode == 'train':
    #####
    # TODO: Implement training phase forward pass for inverted dropout.
    # Store the dropout mask in the mask variable.
    #####
    # Replace "pass" statement with your code
    p=1-p
    mask=(torch.rand(*x.shape, device=x.device)<p)/p
    out=x*mask
    #####
    #                                     END OF YOUR CODE
    #####
elif mode == 'test':
    #####
    # TODO: Implement the test phase forward pass for inverted dropout.
    #####
    # Replace "pass" statement with your code
    out=x
    #####
    #                                     END OF YOUR CODE
    #####
```

Dropout backward 在訓練的時候，看哪裡有被 dropout 就不要 pass 過去，沒被 dropout 就還原原來的值。在測試的時候不需要 dropout，值直接 pass 過去。程式碼如下：

```

if mode == 'train':
    #####
    # TODO: Implement
    #####
    # Replace "pass"
    dx=dout*mask
    #####
    #
    #####
elif mode == 'test':
    dx = dout

```

四. Fully connected net:

這裡我只列出我在 loss 修改的部分。

在 forward 時，如果有使用到 dropout 的話，就用 dropout.forward 去看神經元有沒有被關掉，並用 cache 儲存結果，程式碼如下：

```

if self.use_dropout:
    out, cache=Dropout.forward(out, self.dropout_param)
    caches_drop.append(cache)

```

剛剛在 forward 時是先 forward 再看有沒有 dropout，在 backward 時順序就應該反過來，並用 cache.pop() 傳遞中間參數。程式碼如下：

```

if self.use_dropout:
    dout=Dropout.backward(dout, caches_drop.pop())

```

完整的程式碼如下：

```

out=X
caches=[]
caches_drop=[]
for i in range(self.num_layers-1):
    out, cache=Linear_ReLU.forward(out, self.params['W'+repr(i+1)], self.params['b'+repr(i+1)])
    caches.append(cache)
    if self.use_dropout:
        out, cache=Dropout.forward(out, self.dropout_param)
        caches_drop.append(cache)
#####
#                               END OF YOUR CODE
#####
scores, cache = Linear.forward(out, self.params['W'+repr(self.num_layers)], self.params['b'+repr(self.num_layers)])
caches.append(cache)

```

```

data_loss, dscore=softmax_loss(scores, y)
reg_loss=0.0
for i in range(self.num_layers):
    w=self.params['W'+repr(i+1)]
    reg_loss+=self.reg*torch.sum(w**2)*0.5
loss=data_loss+reg_loss
dout, dw, db=Linear.backward(dscore, caches.pop())
grads['W'+repr(self.num_layers)]=dw+self.reg*self.params['W'+repr(self.num_layers)]
grads['b'+repr(self.num_layers)]=db
for i in range(self.num_layers-2, -1, -1):
    if self.use_dropout:
        dout=Dropout.backward(dout, caches_drop.pop())
    dout, dw, db=Linear_ReLU.backward(dout, caches.pop())
    grads['W'+repr(i+1)]=dw+self.reg*self.params['W'+repr(i+1)]
    grads['b'+repr(i+1)]=db

```

貳. 執行結果：

1. 如下圖所示，用 rms prop 算出來的 w 和 cache 裡的參數跟實際上沒差

很多：

```

expected_next_w = torch.tensor([
  [-0.39223849, -0.34037513, -0.28849239, -0.23659121, -0.18467247],
  [-0.132737, -0.08078555, -0.02881884, 0.02316247, 0.07515774],
  [ 0.12716641, 0.17918792, 0.23122175, 0.28326742, 0.33532447],
  [ 0.38739248, 0.43947102, 0.49155973, 0.54365823, 0.59576619]],
  dtype=torch.float64, device='cuda')
expected_cache = torch.tensor([
  [ 0.5976, 0.6126277, 0.6277108, 0.64284931, 0.65804321],
  [ 0.67329252, 0.68859723, 0.70395734, 0.71937285, 0.73484377],
  [ 0.75037008, 0.7659518, 0.78158392, 0.79728144, 0.81302936],
  [ 0.82883269, 0.84469141, 0.86060554, 0.87657507, 0.8926]],
  dtype=torch.float64, device='cuda')

print('next_w error: ', usefuns.grad.rel_error(expected_next_w, next_w))
print('cache error: ', usefuns.grad.rel_error(expected_cache, config['cache']))

next_w error: 4.064797880829826e-09
cache error: 1.8620321382570356e-09

```

2. 如下圖所示，用 adam 算出來的 w 和 Mean Squared Gradient 和 Root Mean Squared Gradient 跟實際上沒差很多：

```
expected_next_w = torch.tensor([
  [-0.40094747, -0.34836187, -0.29577703, -0.24319299, -0.19060977],
  [-0.1380274, -0.08544591, -0.03286534, 0.01971428, 0.0722929],
  [ 0.1248705, 0.17744702, 0.23002243, 0.28259667, 0.33516969],
  [ 0.38774145, 0.44031188, 0.49288093, 0.54544852, 0.59801459]],
  dtype=torch.float64, device='cuda')
expected_v = torch.tensor([
  [ 0.69966, 0.68908382, 0.67851319, 0.66794809, 0.65738853],
  [ 0.64683452, 0.63628604, 0.6257431, 0.61520571, 0.60467385],
  [ 0.59414753, 0.58362676, 0.57311152, 0.56260183, 0.55209767],
  [ 0.54159906, 0.53110598, 0.52061845, 0.51013645, 0.49966,
  ]],
  dtype=torch.float64, device='cuda')
expected_m = torch.tensor([
  [ 0.48, 0.49947368, 0.51894737, 0.53842105, 0.55789474],
  [ 0.57736842, 0.59684211, 0.61631579, 0.63578947, 0.65526316],
  [ 0.67473684, 0.69421053, 0.71368421, 0.73315789, 0.75263158],
  [ 0.77210526, 0.79157895, 0.81105263, 0.83052632, 0.85
  ]],
  dtype=torch.float64, device='cuda')

# You should see relative errors around e-7 or less
print('next_w error: ', usefuns.grad.rel_error(expected_next_w, next_w))
print('v error: ', usefuns.grad.rel_error(expected_v, config['v']))
print('m error: ', usefuns.grad.rel_error(expected_m, config['m']))

next_w error: 3.756728297598868e-09
v error: 3.4048987160545265e-09
m error: 2.786377729853651e-09
```



3. 如下圖所示，比較 sgd 和 sgd_momentum 和 adam 和 rms prop 之間的準確率，發現 adam>rms prop>sgd momentum>sgd，loss 則是 adam<rms prop<sgd momentum<sgd 並把結果 plot 出來。

```

running with adam
(Time 0.00 sec; Iteration 1 / 200) loss: 2.302573
(Epoch 0 / 5) train acc: 0.119000; val_acc: 0.121700
(Epoch 1 / 5) train acc: 0.281000; val_acc: 0.244400
(Epoch 2 / 5) train acc: 0.353000; val_acc: 0.313900
(Epoch 3 / 5) train acc: 0.415000; val_acc: 0.347500
(Epoch 4 / 5) train acc: 0.405000; val_acc: 0.354800
(Epoch 5 / 5) train acc: 0.453000; val_acc: 0.367100

running with rmsprop
(Time 0.00 sec; Iteration 1 / 200) loss: 2.303421
(Epoch 0 / 5) train acc: 0.125000; val_acc: 0.113100
(Epoch 1 / 5) train acc: 0.216000; val_acc: 0.222000
(Epoch 2 / 5) train acc: 0.278000; val_acc: 0.266400
(Epoch 3 / 5) train acc: 0.283000; val_acc: 0.280600
(Epoch 4 / 5) train acc: 0.338000; val_acc: 0.287400
(Epoch 5 / 5) train acc: 0.313000; val_acc: 0.288200

```



4. 如下圖所示，輸出的平均值在訓練和測試期間大致相同。在訓練期間，設為零的輸出數應大致等於 p 。而在測試期間，不應有任何輸出被設為零。

```
Running tests with p = 0.25
Mean of input: 9.997330335850453
Mean of train-time output: 10.008798328321282
Mean of test-time output: 9.997330335850453
Fraction of train-time output set to zero: 0.2490839958190918
Fraction of test-time output set to zero: 0.0

Running tests with p = 0.4
Mean of input: 9.997330335850453
Mean of train-time output: 9.977839345169954
Mean of test-time output: 9.997330335850453
Fraction of train-time output set to zero: 0.40116000175476074
Fraction of test-time output set to zero: 0.0

Running tests with p = 0.7
Mean of input: 9.997330335850453
Mean of train-time output: 10.006132149340187
Mean of test-time output: 9.997330335850453
Fraction of train-time output set to zero: 0.6999160051345825
Fraction of test-time output set to zero: 0.0
```

5. 如下圖所示，dropout backward 出來的梯度與實際上的誤差很小。

```
dropout_param = {'mode': 'train', 'p': 0.2, 'seed': 0}
out, cache = Dropout.forward(x, dropout_param)
dx = Dropout.backward(dout, cache)
dx_num = usefuns.grad.compute_numeric_gradient(lambda xx: Dropout.forward(xx, dropout_param)[0], x, dout)

# Error should be around e-10 or less
print('dx relative error: ', usefuns.grad.rel_error(dx, dx_num))

dx relative error: 3.914942325636866e-09
```

6. 如下圖所示，在 fully connect net 加上 dropout，發現跟實際計算出來的數值誤差很小。

```
Running check with dropout = 0
Initial loss: 2.3053575717037686
W1 relative error: 6.06e-08
W2 relative error: 1.02e-07
W3 relative error: 5.89e-08
b1 relative error: 1.28e-07
b2 relative error: 2.05e-08
b3 relative error: 3.41e-09

Running check with dropout = 0.25
Initial loss: 2.3149601124989854
W1 relative error: 3.64e-08
W2 relative error: 4.60e-08
W3 relative error: 3.14e-08
b1 relative error: 6.83e-08
b2 relative error: 3.17e-08
b3 relative error: 3.63e-09

Running check with dropout = 0.5
Initial loss: 2.3021646399354614
W1 relative error: 7.21e-08
W2 relative error: 2.93e-08
W3 relative error: 2.56e-08
b1 relative error: 5.59e-08
b2 relative error: 2.64e-08
b3 relative error: 2.76e-09
```

6. 如下圖所示，比較不同的 dropout 和 hidden size，發現沒有用 dropout 時 train 的準確率高但 validation 的準確率低。用 dropout 時，train 的準確率雖然比沒有用 dropout 低，但 validation 的準確率卻比沒有用 dropout 高，可以降低 overfitting。

```
Training a model with dropout=0.00 and width=256
(Time 0.01 sec; Iteration 1 / 3900) loss: 2.304467
(Epoch 0 / 100) train acc: 0.193000; val_acc: 0.198200
(Epoch 10 / 100) train acc: 0.742000; val_acc: 0.482600
(Epoch 20 / 100) train acc: 0.876000; val_acc: 0.474400
(Epoch 30 / 100) train acc: 0.913000; val_acc: 0.467000
(Epoch 40 / 100) train acc: 0.951000; val_acc: 0.459700
(Epoch 50 / 100) train acc: 0.973000; val_acc: 0.462600
(Epoch 60 / 100) train acc: 0.930000; val_acc: 0.458900
(Epoch 70 / 100) train acc: 0.989000; val_acc: 0.469900
(Epoch 80 / 100) train acc: 1.000000; val_acc: 0.481900
(Epoch 90 / 100) train acc: 1.000000; val_acc: 0.483900
(Epoch 100 / 100) train acc: 1.000000; val_acc: 0.481300
```

```
Training a model with dropout=0.00 and width=512
(Time 0.00 sec; Iteration 1 / 3900) loss: 2.302387
(Epoch 0 / 100) train acc: 0.239000; val_acc: 0.220000
(Epoch 10 / 100) train acc: 0.723000; val_acc: 0.484900
(Epoch 20 / 100) train acc: 0.891000; val_acc: 0.470500
(Epoch 30 / 100) train acc: 0.951000; val_acc: 0.481100
(Epoch 40 / 100) train acc: 0.944000; val_acc: 0.475000
(Epoch 50 / 100) train acc: 0.937000; val_acc: 0.472700
(Epoch 60 / 100) train acc: 0.958000; val_acc: 0.479500
(Epoch 70 / 100) train acc: 0.937000; val_acc: 0.463900
(Epoch 80 / 100) train acc: 0.969000; val_acc: 0.470400
(Epoch 90 / 100) train acc: 0.979000; val_acc: 0.477900
(Epoch 100 / 100) train acc: 0.954000; val_acc: 0.468400
```

```
Training a model with dropout=0.50 and width=512
(Time 0.00 sec; Iteration 1 / 3900) loss: 2.306579
(Epoch 0 / 100) train acc: 0.233000; val_acc: 0.227000
(Epoch 10 / 100) train acc: 0.603000; val_acc: 0.467600
(Epoch 20 / 100) train acc: 0.666000; val_acc: 0.485400
(Epoch 30 / 100) train acc: 0.714000; val_acc: 0.489700
(Epoch 40 / 100) train acc: 0.807000; val_acc: 0.492400
(Epoch 50 / 100) train acc: 0.847000; val_acc: 0.497000
(Epoch 60 / 100) train acc: 0.873000; val_acc: 0.497500
(Epoch 70 / 100) train acc: 0.896000; val_acc: 0.502000
(Epoch 80 / 100) train acc: 0.913000; val_acc: 0.487300
(Epoch 90 / 100) train acc: 0.936000; val_acc: 0.491300
(Epoch 100 / 100) train acc: 0.923000; val_acc: 0.493500
```

